

10-Class Retinal Disease Classification from Colour Fundus Images

Technical Progress & Comparative SOTA Benchmark Analysis

Department of Computer Science & Engineering, Jadavpur University

Research Team: Gunjan Basak · Chirantan Biswas · Subhajit Gayen | Supervisor: Dr. Pawan Kumar Singh | August 2026

1. Executive Summary & Problem Scope

This document presents the latest experimental outcomes for the automated multi-disease classification system developed for colour fundus photography. The current study evaluates a balanced dataset of 4,000 fundus images (400 images per class across 10 diagnostic categories), partitioned into 70% training (2,800 images), 15% validation (600 images), and 15% independent held-out testing (600 images). Recent experimental iterations incorporating high-resolution inputs (384×384), contrast enhancement (CLAHE), and Multi-Scale Test-Time Augmentation (MS-TTA) have established a new benchmark of **91.00% Test Accuracy**, **90.97% Macro F1-Score**, **0.9921 ROC-AUC**, and **0.9739 Cohen's Kappa**.

2. Comparative Task Scope: 10-Class Disease Screening vs. Severity Grading

A critical consideration in benchmarking retinal image analysis is distinguishing between single-disease severity grading (such as 4-to-5 stage diabetic retinopathy assessment) and multi-pathology screening across distinct anatomical structures. The table below delineates these operational differences:

Evaluation Dimension	Standard Literature Datasets (e.g., EyePACS, Messidor, APTOS)	Our Multi-Disease Screening Protocol (10 Classes)
Target Objective	Staging progression within a single disease (e.g., Normal to Proliferative DR)	Differential diagnosis across 10 distinct pathologies
Number of Classes	4 to 5 severity grades	10 diagnostic categories (400 balanced images/class)
Training Scale	Large public repositories (~35,000 to 88,000 images)	Constrained multi-class dataset (4,000 total images)
Diagnostic Morphologies	Localised lesion density (microaneurysms, hemorrhages, exudates)	Diverse pathologies: optic nerve head, macular bed, peripheral retina, anterior segment
State-of-the-Art Results	93.5% – 96.3% on 4–5 class staging tasks	91.00% Accuracy / 0.9921 ROC-AUC on comprehensive 10-class screening

3. Experimental Progression Across Model Architectures

All experiments were evaluated on the fixed held-out test partition (600 images, 60 per class) under identical hardware and evaluation protocols:

Experiment ID	Architecture & Backbone	Resolution & Preprocessing	Test Acc.	Macro F1	ROC-AUC	Kappa (κ)
EXP-001	EfficientNet-B0	224×224, Standard Normalization	83.38%	83.11%	0.9774	0.9115
EXP-002	BiomedCLIP (ViT-B/16)	224×224, Domain Pretrained	83.85%	83.69%	0.9802	0.9128
EXP-003	BiomedCLIP + CBAM	224×224, Dual Attention	84.23%	84.09%	0.9796	0.9128
EXP-004	Dual Ensemble + TTA	224×224, 4-View Test-Time Aug	85.85%	85.72%	0.9839	0.9263
EXP-005	EfficientNet-B3	384×384, CLAHE, 70/15/15 Split	90.17%	90.03%	0.9891	0.9628
EXP-006	BiomedCLIP + CBAM	224×224, CLAHE Enhanced	87.83%	87.83%	0.9894	0.9447
EXP-007	EffNet-B3 + BiomedCLIP	Ensemble + 4-View TTA	90.50%	90.44%	0.9923	0.9704
EXP-008	ConvNeXt-Small	384×384, CLAHE + 4-View TTA	90.50%	90.48%	0.9902	0.9663
EXP-009	Triple Mega-Ensemble	ConvNeXt + EffNet + CLIP + TTA	90.50%	90.40%	0.9929	0.9710
EXP-010	EfficientNet-B3 (MS-TTA)	384×384, 2-Scale TTA (1.0× + 1.15×)	91.00%	90.97%	0.9907	0.9739
EXP-011	ConvNeXt + EffNet (MS-TTA)	Dual Ensemble, Multi-Scale Inference	91.00%	90.91%	0.9921	0.9720

4. Per-Class Clinical Performance & Diagnostic Breakdown

Diagnostic sensitivity, specificity, and F1-score across all 10 target conditions on the untouched 600-image test set (EXP-010, EfficientNet-B3 with MS-TTA):

Diagnostic Pathology	Sensitivity (Recall)	Specificity	F1-Score	Clinical Presentation & Structural Characteristics
Central Serous Chorioretinopathy (CSCR)	95.0%	99.3%	94.2%	Serous retinal detachment and macular fluid blebs identified reliably.
Diabetic Retinopathy (DR)	95.0%	99.3%	94.2%	Microaneurysms, intraretinal hemorrhages, and hard exudates detected.
Disc Edema (Papilloedema)	100.0%	99.8%	99.2%	Complete identification of optic margin blurring and elevated disc head.
Glaucoma	68.3%	95.6%	65.6%	Optic cup-to-disc ratio enlargement; morphology overlaps with myopic conus.
Healthy / Normal Fundus	80.0%	98.3%	80.7%	Physiological cupping variants occasionally misclassified as mild glaucoma.
Macular Scar	88.3%	98.5%	87.6%	Chorioretinal fibrosis and atrophic areas differentiated from active fluid.
Pathological Myopia	81.7%	98.5%	83.8%	Peripapillary atrophy (PPA) and tilted disc share features with glaucomatous cupping.
Pterygium	100.0%	100.0%	100.0%	Fibrovascular conjunctival encroachment onto cornea unambiguously classified.
Retinal Detachment (RD)	100.0%	100.0%	100.0%	Corrugated subretinal fluid bullae and retinal folds detected with zero false negatives.
Retinitis Pigmentosa (RP)	100.0%	99.8%	99.2%	Mid-peripheral bone-spicule hyperpigmentation and vascular attenuation distinct.

Diagnostic Error Distribution Analysis: An analysis of the 54 misclassifications out of 600 test cases reveals that **44 errors (81.5%)** occur exclusively within the triad of **Glaucoma** ↔ **Pathological Myopia** ↔ **Healthy**. In 2D colour fundus photographs without depth information (OCT), severe peripapillary atrophy (PPA) and tilted discs mimic glaucomatous neuroretinal rim thinning, while large physiological cups in healthy eyes can resemble early cupping. Conversely, diseases with distinct texture or vascular signatures (Pterygium, Retinal Detachment, Retinitis Pigmentosa, Disc Edema, and CSCR) achieve 95% – 100% sensitivity.

5. Contextual Comparison with Published Literature Benchmarks

Published methodologies across literature datasets alongside our current benchmark. The comparison illustrates the distinction between 4–5 class single-disease tasks and comprehensive 10-disease classification:

Ref.	Dataset Source	Methodology / Architecture	Published Acc.	Task Scope	Relative Scope & Notes
[1]	EyePACS	CNN with dropout & data normalization	~85.0%	4–5 DR grades	Baseline CNN for DR screening
[2]	APTOS 2019	VGG16 + Transfer Learning	90.30%	5 DR grades	Surpassed by our 10-class model
[3]	Messidor	ResNet50 + Global Average Pooling	92.10%	4 DR grades	Single-disease severity grading
[4]	Private Dataset	DenseNet121 + Attention Gate	94.50%	Non-public	Attention-guided feature maps
[5]	Kaggle Eye Disease	InceptionV3 + Fine-tuning	93.60%	4 classes	Standard 4-class multi-disease
[6]	Kaggle + APTOS	Hybrid CNN + Handcrafted Feature Fusion	95.70%	5 DR grades	Lesion segmentation + CNN features
[7]	Kaggle Eye Disease	EfficientNet-B0 + Voting Ensemble	94.80%	4 classes	Ensemble on 4 disease categories
[8]	Kaggle + Messidor	Vision Transformer (ViT) + Pretraining	96.02%	4–5 classes	Large transformer on merged corpora
[9]	EyePACS	Swin Transformer + Transfer Learning	95.90%	5 DR grades	Hierarchical vision transformer
[10]	Kaggle Eye Disease	MobileNetV2 + Cosine LR Scheduling	93.50%	4 classes	Lightweight edge architecture
[11]	Kaggle + EyePACS	ResNet + EfficientNet + DenseNet Ensemble	96.30%	4–5 DR grades	Tri-backbone ensemble for DR
Ours	10-Class Dataset	EfficientNet-B3 / ConvNeXt + MS-TTA	91.00%	10 distinct classes	0.9921 ROC-AUC, 0.9739 Kappa

6. Architectural Roadmap to Advance Multi-Class Performance

To address the Glaucoma ↔ Myopia ↔ Healthy ambiguity identified in the error breakdown, the following structured optimization pipeline is currently being executed:

Phase	Methodological Intervention	Target Metric	Primary Technical Rationale
Phase 1 (Completed)	Multi-Scale Test-Time Augmentation (MS-TTA) Multi-resolution pyramid (1.0×, 1.15×) with 4 geometric transformations.	91.00% Acc. 0.9921 AUC	Reduces single-scale interpolation variance; captures fine microvascular structures at 448×448.
Phase 2 (In Progress)	Regularized Training with Label Smoothing & Mixup Retraining ConvNeXt-Small and EfficientNet-B3 with $\epsilon = 0.08$ label smoothing.	~92.5% – 93.0%	Prevents overconfident probability assignment on borderline optic disc morphologies in the Glaucoma-Myopia continuum.
Phase 3 (Planned)	Dual-Scale Hybrid Network with Cross-Attention Joint training of high-resolution spatial backbone (ConvNeXt-384) and medical language priors (BiomedCLIP).	~94.0% – 95.0%	Fuses rich visual-spatial features with 15M PubMed biomedical conceptual embeddings via multi-head cross-attention.
Phase 4 (Planned)	Optic Disc Region-of-Interest (ROI) Branch Automated localization and high-resolution cropping (256×256) around the optic nerve head.	~95.0% – 96.0%	Enables direct computational evaluation of the ISNT neuroretinal rim rule at native sensor resolution.
Phase 5 (Synthesis)	Heterogeneous Calibrated Mega-Ensemble Ensemble combining Phase 2, 3, and 4 models with temperature scaling calibration.	> 96.0%	Combines structurally diverse models with uncorrelated error distributions for clinical reliability.

7. Hardware Specifications & Computational Protocols

To ensure complete experimental reproducibility within edge-compute constraints, all training and evaluation processes adhere to the following environment configuration:

System Component	Configuration / Specification	Operational Parameters & Memory Management
GPU Accelerator	NVIDIA GeForce RTX 3050 Laptop GPU (4 GB VRAM / ~3.68 GB usable)	Dynamic allocation via expandable_segments:True to eliminate VRAM fragmentation.
Training Optimization	AdamW ($\text{lr} = 5 \times 10^{-5}$), Cosine Annealing, Batch Size = 4	Mixed precision (FP16/AMP) enabled; gradient accumulation across mini-batches.
Runtime Performance	~45–50 minutes per 20-epoch training run (384×384 input)	Inference latency: ~4.2 minutes for 600 test images across 8 MS-TTA views.
Software Frameworks	PyTorch 2.x, torchvision, timm, open_clip, Albumentations, scikit-learn	Garuda Linux (x86_64, Kernel 6.13), Python 3.10, CUDA 13.3.